

13-网络传输机制(TCP)实验二报告

陈彦帆 2018K8009918002

1、实验内容

(1) 实现 TCP 数据传输：将数据封装到数据包并发送；收到数据和 ACK 时的相应处理。

(2) 实现流量控制：通过调整 `recv_window` 来表达自己的接收能力。

(3) 实现 `tcp_sock` 相关函数：`read`, `write`。

(4) 在节点 h1 上运行 TCP server，在 h2 上运行 TCP client，向 h1 发送字符串，h2 echo，h1 打印收到的信息。然后，在一端用 `tcp_stack.py` 替换 `tcp_stack` 执行，测试另一端。

(5) 修改 `tcp_apps.c`。在节点 h1 上运行 TCP server，在 h2 上运行 TCP client，向 h1 发送一个大小为几 MB 的文件，比较 h1 接收的文件与 h2 发送的文件是否相同。然后，在一端用修改后的 `tcp_stack.py` 替换 `tcp_stack` 执行，测试另一端。

2、实验流程

(1) 完成代码编写并编译。

(2) 测试正确性。

(实验内容(4)和(5)的测试流程相同，后者在测试前需要执行 `bash create_randfile.sh` 在测试后需要执行 `diff server-output.dat client-input.dat`)

首先通过 mininet 环境(`tcp_topo.py`)构建网络拓扑如下图：



图 1

在节点 h1 上运行 TCP server，在 h2 上运行 TCP client，进行抓包测试：

执行：

```
sudo python tcp_topo.py
```

```
mininet> xterm h1 h2
```

```
h1# wireshark
```

```
h1# ./tcp_stack server 10001
```

```

or h1# python tcp_stack.py server 10001
h2# ./tcp_stack client 10.0.0.1 10001
or h2# python tcp_stack.py client 10.0.0.1 10001
mininet> quit

```

3、分析与实现

(1) 实现 TCP 数据传输。

① 收到数据和 ACK 时的相应处理：

首先检查收到的包是否有效。收到的包的字节序必须在本端的接收窗口以内。在本实验中，不能处理乱序到达的包，这些包将被丢弃并报错。

每收到一个有效的包，更新对应的 sock 信息：更新 rcv_next 为收到数据的第一个字节加上数据长度；如果收到的包包含 ACK，更新 snd_una；根据收到包的接收窗口更新 sock 的发送窗口；如果收到的包有新的信息，发送 ACK 包。

```

static inline int update_tsk(struct tcp_sock *tsk, struct tcp_cb *cb)
{
    tsk->rcv_nxt = cb->seq_end;
    if(cb->flags & TCP_ACK)
        tsk->snd_una = max(tsk->snd_una,cb->ack);
    tcp_update_window_safe(tsk,cb);
    if(tsk->state != TCP_LISTEN && less_than_32b(cb->seq,cb->seq_end))
        tcp_send_control_packet(tsk,TCP_ACK);
    return 0;
}

```

如果收到的包包含数据部分，则将数据部分写入环形缓存。出于简单考虑，若缓存放不下整个包，整个包将被丢弃。

环形缓存是一个生产者-消费者模型，若缓存为空，消费者被挂起。但若缓存为满，生产者不需要被挂起，而是把超出缓存的部分丢弃。

在访问缓存时，需要上锁。这里用的是 tsk->wait_rcv->lock。当一方被挂起时，调用 pthread_cond_wait，将锁交出并等待唤醒。唤醒时，将尝试获取锁。

```

void handle_tcp_data(struct tcp_sock *tsk, struct tcp_cb *cb)
{
    if(!cb->pl_len) return;
    pthread_mutex_lock(&tsk->wait_rcv->lock);
    if(ring_buffer_free(tsk->rcv_buf) < cb->pl_len)
        log(DEBUG,"RCV BUFFER FULL. DROPED.");
    else{

```

```

    write_ring_buffer(tsk->rcv_buf,cb->payload,cb->pl_len);
    wake_with_lock(tsk->wait_recv);
}
pthread_mutex_unlock(&tsk->wait_recv->lock);
}

```

② 将数据封装到数据包并发送

框架已经实现。

(2) 实现流量控制：通过调整 `recv_window` 来表达自己的接收能力。

每收到一个数据包，调整当前 sock 的发送窗口为对端的接收窗口大小。调整接收窗口为环形缓存的剩余大小。

(3) 实现 `tcp_sock` 相关函数：

① read

函数调用者是消费者。若环形缓存为空，则交出锁进入等待。否则读取数据并返回。若唤醒时发现 sock 进入 `CLOSE_WAIT` 状态，说明发送方断开连接，返回 0。

```

int tcp_sock_read(struct tcp_sock *tsk, char *buf, int len)
{
    pthread_mutex_lock(&tsk->wait_recv->lock);
    while(ring_buffer_empty(tsk->rcv_buf)){
        if(tsk->state == TCP_CLOSE_WAIT)
            return 0;
        sleep_with_lock(tsk->wait_recv);
    }
    int readlen = read_ring_buffer(tsk->rcv_buf,buf,len);
    tsk->rcv_wnd = ring_buffer_free(tsk->rcv_buf);
    pthread_mutex_unlock(&tsk->wait_recv->lock);
    return readlen;
}

```

② write

每次发送的大小为 $\min(\text{MSS}, \text{data_left}, \text{send_window})$ 。若发送窗口为 0，则交出锁等待。否则调用 `tcp_send_packet` 进行发送。

```

int tcp_sock_write(struct tcp_sock *tsk, char *buf, int len)
{
    int tot = len;
    while(len>0){
        int headerlen = ETHER_HDR_SIZE + IP_BASE_HDR_SIZE + TCP_BASE_HDR_SIZE;
        int sendlen = min(len,ETH_FRAME_LEN-headerlen);
        pthread_mutex_lock(&tsk->wait_send->lock);
        while(tsk->snd_wnd==0)
            sleep_with_lock(tsk->wait_send);

```

```

    sendlen = min(sendlen,tsk->snd_wnd);
    pthread_mutex_unlock(&tsk->wait_send->lock);
    int pkt_len = sendlen + headerlen;
    char *pkt = malloc(pkt_len);
    memcpy(pkt+headerlen,buf,sendlen);
    tcp_send_packet(tsk,pkt,pkt_len);
    buf += sendlen, len -= sendlen;
}
return tot;
}

```

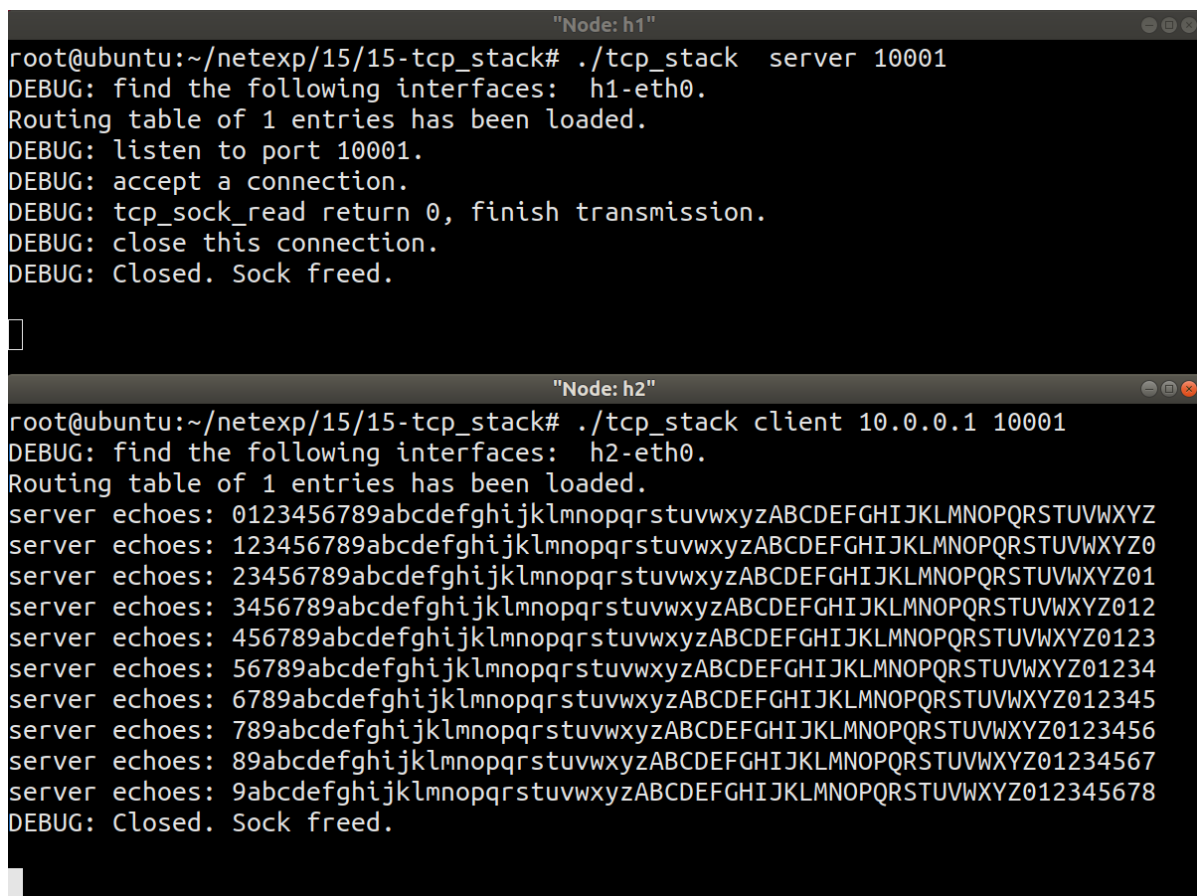
4、实验结果

网络拓扑如图 1。

(1) 测试 1：字符串回显。

shell 输出结果如图 2。抓包结果如图 3。将 tcp_stack 一端替换为 python 程序后，结果如图 4-5。

经过比对，我实现的 tcp 读写功能正确，任意一端替换为 python 程序后，结果也正确。



```

"Node: h1"
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack server 10001
DEBUG: find the following interfaces: h1-eth0.
Routing table of 1 entries has been loaded.
DEBUG: listen to port 10001.
DEBUG: accept a connection.
DEBUG: tcp_sock_read return 0, finish transmission.
DEBUG: close this connection.
DEBUG: Closed. Sock freed.

"Node: h2"
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack client 10.0.0.1 10001
DEBUG: find the following interfaces: h2-eth0.
Routing table of 1 entries has been loaded.
server echoes: 0123456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ
server echoes: 123456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0
server echoes: 23456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ01
server echoes: 3456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ012
server echoes: 456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0123
server echoes: 56789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ01234
server echoes: 6789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ012345
server echoes: 789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0123456
server echoes: 89abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ01234567
server echoes: 9abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ012345678
DEBUG: Closed. Sock freed.

```

图 2

4	0.021951368	10.0.0.2	10.0.0.1	TCP	54	12345 → 10001	[SYN]	Seq=0 Win=65535 Len=0
5	0.032907658	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[SYN, ACK]	Seq=0 Ack=1 Win=65535 Len=0
6	0.043883450	10.0.0.2	10.0.0.1	TCP	54	12345 → 10001	[ACK]	Seq=1 Ack=1 Win=65535 Len=0
7	0.043885192	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=1 Ack=1 Win=65535 Len=62
8	0.054845141	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=1 Ack=63 Win=65535 Len=77
9	1.066704684	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=63 Ack=78 Win=65535 Len=62
10	1.077696892	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=78 Ack=125 Win=65535 Len=77
11	2.089710992	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=125 Ack=155 Win=65535 Len=62
12	2.100478025	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=155 Ack=187 Win=65535 Len=77
13	3.110728332	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=187 Ack=232 Win=65535 Len=62
14	3.120835932	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=232 Ack=249 Win=65535 Len=77
15	4.132093461	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=249 Ack=309 Win=65535 Len=62
16	4.143068089	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=309 Ack=311 Win=65535 Len=77
17	5.154016642	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=311 Ack=386 Win=65535 Len=62
18	5.165087926	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=386 Ack=373 Win=65535 Len=77
19	6.176923605	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=373 Ack=463 Win=65535 Len=62
20	6.187813686	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=463 Ack=435 Win=65535 Len=77
21	7.199692024	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=435 Ack=540 Win=65535 Len=62
22	7.210602439	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=540 Ack=497 Win=65535 Len=77
23	8.222009123	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=497 Ack=617 Win=65535 Len=62
24	8.232909159	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=617 Ack=559 Win=65535 Len=77
25	9.244258899	10.0.0.2	10.0.0.1	TCP	116	12345 → 10001	[PSH, ACK]	Seq=559 Ack=694 Win=65535 Len=62
26	9.255215116	10.0.0.1	10.0.0.2	TCP	131	10001 → 12345	[PSH, ACK]	Seq=694 Ack=621 Win=65535 Len=77
27	10.266593180	10.0.0.2	10.0.0.1	TCP	54	12345 → 10001	[FIN, ACK]	Seq=621 Ack=771 Win=65535 Len=0
28	10.277027908	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[ACK]	Seq=771 Ack=622 Win=65535 Len=0
29	10.277041935	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[FIN, ACK]	Seq=771 Ack=622 Win=65535 Len=0
30	10.287128586	10.0.0.2	10.0.0.1	TCP	54	12345 → 10001	[ACK]	Seq=622 Ack=772 Win=65535 Len=0

图 3

```

"Node: h1"
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack server 10001
DEBUG: find the following interfaces: h1-eth0.
Routing table of 1 entries has been loaded.
DEBUG: listen to port 10001.
DEBUG: accept a connection.
DEBUG: tcp_sock_read return 0, finish transmission.
DEBUG: close this connection.
DEBUG: Closed. Sock freed.

"Node: h2"
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack.py client 10.0.0.1 10001
server echoes: 0123456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ0
server echoes: 123456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ01
server echoes: 23456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ012
server echoes: 3456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ0123
server echoes: 456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ01234
server echoes: 56789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ012345
server echoes: 6789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ0123456
server echoes: 789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ01234567
server echoes: 89abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ012345678
server echoes: 9abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZKLMNOPQRSTUVWXYZ0123456789
root@ubuntu:~/netexp/15/15-tcp_stack#

```

图 4 client python

```
"Node: h1"
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack.py server 10001
('10.0.0.2', 12345)
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
<type 'str'>
root@ubuntu:~/netexp/15/15-tcp_stack#

"Node: h2"
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack client 10.0.0.1 10001
DEBUG: find the following interfaces: h2-eth0.
Routing table of 1 entries has been loaded.
server echoes: 0123456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ
server echoes: 123456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0
server echoes: 23456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ01
server echoes: 3456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ012
server echoes: 456789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0123
server echoes: 56789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ01234
server echoes: 6789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ012345
server echoes: 789abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0123456
server echoes: 89abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ01234567
server echoes: 9abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ012345678
DEBUG: Closed. Sock freed.
```

图 5 server python

(2) 测试 2: 上传大文件 (约 4MB)

shell 输出结果如图 6。传输后, 通过 diff 命令比对, 接收方收到的文件与发送方上传的完全相同。将 tcp_stack 一端替换为 python 程序后, 结果如图 7-8。同样通过了 diff 比对。

```
"Node: h2"
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack client 10.0.0.1 10001
DEBUG: find the following interfaces: h2-eth0.
Routing table of 1 entries has been loaded.
DEBUG: Closed. Sock freed.

"Node: h1"
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack server 10001
DEBUG: find the following interfaces: h1-eth0.
Routing table of 1 entries has been loaded.
DEBUG: listen to port 10001.
DEBUG: accept a connection.
tot:4052632
DEBUG: close this connection.
DEBUG: Closed. Sock freed.

alphabet@ubuntu:~/netexp/15/15-tcp_stack$ diff server-output.dat client-input.dat
alphabet@ubuntu:~/netexp/15/15-tcp_stack$
```

图 6


```
"Node: h1"
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py server 10001
('10.0.0.2', 39102)
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py server 10001
('10.0.0.2', 39106)
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py server 10001
('10.0.0.2', 39112)
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py server 10001
('10.0.0.2', 12345)
root@ubuntu:~/netexp/15/15-tcp_stack# █

"Node: h2"
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py client 10.0.0.1 1000
1
Traceback (most recent call last):
  File "tcp_stack2.py", line 52, in <module>
    client(sys.argv[2], sys.argv[3])
  File "tcp_stack2.py", line 42, in client
    s.write(data)
AttributeError: '_socketobject' object has no attribute 'write'
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py client 10.0.0.1 1000
1
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py client 10.0.0.1 1000
1
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack client 10.0.0.1 10001
DEBUG: find the following interfaces: h2-eth0.
Routing table of 1 entries has been loaded.
DEBUG: Closed. Sock freed.
█
```

图 7 python server

```
root@ubuntu:~/netexp/15/15-tcp_stack# python tcp_stack2.py client 10.0.0.1 1000
1
root@ubuntu:~/netexp/15/15-tcp_stack# █
root@ubuntu:~/netexp/15/15-tcp_stack# ./tcp_stack server 10001
DEBUG: find the following interfaces: h1-eth0.
Routing table of 1 entries has been loaded.
DEBUG: listen to port 10001.
DEBUG: accept a connection.
tot:4052632
DEBUG: close this connection.
DEBUG: Closed. Sock freed.
```

图 8 python client

5、遇到的问题

(1) out of order

若发送方发送间隔小于 RT_{prop} ，可能会造成数据包乱序，相应的抓包结果如图 9。本实验无法处理乱序包，故将发送间隔调大，并调小模拟器的 RTT。

3179	0.896151415	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[ACK] Seq=1 Ack=2228801 Win=65535 Len=0
3180	0.896154773	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[ACK] Seq=1 Ack=2230201 Win=65535 Len=0
3181	0.896155264	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[ACK] Seq=1 Ack=2231601 Win=65535 Len=0
3182	0.897099925	10.0.0.1	10.0.0.2	TCP	1454	[TCP Previous segment not captured] 12345 → 10001 [PSH, ACK] Seq=2233001 Ack=1 Win=65535 Len=1400	
3183	0.897093642	10.0.0.2	10.0.0.1	TCP	1454	[TCP Retransmission] 12345 → 10001 [PSH, ACK] Seq=2231601 Ack=1 Win=65535 Len=1400	
3184	0.897100309	10.0.0.2	10.0.0.1	TCP	1454	12345 → 10001 [PSH, ACK] Seq=2234401 Ack=1 Win=65535 Len=1400	
3185	0.897100543	10.0.0.2	10.0.0.1	TCP	1454	12345 → 10001 [PSH, ACK] Seq=2235801 Ack=1 Win=65535 Len=1400	
3186	0.898120679	10.0.0.2	10.0.0.1	TCP	1454	12345 → 10001 [PSH, ACK] Seq=2237201 Ack=1 Win=65535 Len=1400	
3187	0.898168863	10.0.0.1	10.0.0.2	TCP	54	10001 → 12345	[ACK] Seq=1 Ack=2234401 Win=65535 Len=0

图 9

(2) socket raw send buffer FULL

若发送过快，socket 可能会出现如下提示：

Send raw packet failed: No buffer space available

解决方法：增加发送间隔。

(3) recv buffer FULL

若发送过快，发送方尚未收到接收方的 ACK 包时已经发送了大量数据，来不及调整发送窗口，造成接收方缓存爆满。解决方法：解决方法：增加发送间隔。

(4) 文件写入失败

接收方在向本地写文件完毕后，应调用 `fclose`，否则文件可能写入失败。

(5) ACK 超时

接收方的 ACK 可与数据一并发送。在本实验中，client 间隔 1s 发送一次字符串，此时 ACK 已经超时，python server 会重传。解决办法：收到数据包后马上发送 ACK。

(6) 唤醒条件不满足：

对于框架给出的 `sleep_on` 和 `wake_up` 函数，若先调用了 `wake_up`，则调用 `sleep_on` 时不会睡眠而是立即返回，不满足唤醒条件，造成错误。

解决方法：重写 `sleep_on` 和 `wake_up` 函数，使 `sleep_on` 只有在调用后被唤醒才会返回。